

SIT789 – Robotics, Computer Vision and Speech Processing

Pass Task 6.1: 3D Vision

Objectives

The objectives of this lab include:

- Practising 3D projection in pin-hole models
 - Practising camera pose estimation
 - Practising 3D reconstruction with the structure-from-motion (sfm) approach
-

Tasks

1. 3D projection

In this section, we will implement the camera matrix for a pin-hole camera model and then will perform 3D projection (i.e., projecting a point in 3D space onto a 2D image).

Recall that the camera matrix P of a pin-hole camera can be written as $P = K [R | t]$ (see Eq (2) in slide 6, in Lecture 6 slides). For simplicity, we initialise K , R , and t as follows.

```
import numpy as np

#setup camera with a simple camera matrix P
f = 100
cx = 200
cy = 200
K = np.array([[f, 0, cx], [0, f, cy], [0, 0, 1]])
I = np.eye(3)
t = np.array([[0], [0], [0]])
P = np.dot(K, np.hstack((I, t)))
```

We then define a 3D projection method as,

```
def project(P, X): #X is an array of 3D points
    x = np.dot(P, X)
    for i in range(3): #convert to inhomogeneous coordinates
        x[i] /= x[2]
    return x
```

To test the camera matrix P , we use the 3D point set in [house.p3d](http://www.robots.ox.ac.uk/~vgg/data/mview/) supplied in OnTrack. This dataset is from <http://www.robots.ox.ac.uk/~vgg/data/mview/>.

You first need to download the file [house.p3d](#) into your working directory and read data from the file using the following code.

```
#load data
points_3D = np.loadtxt('house.p3d').T #T means tranpose
points_3D = np.vstack((points_3D, np.ones(points_3D.shape[1])))
```

To visualise the 3D data in [house.p3d](#), we can perform.

```
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import axes3d

fig = plt.figure(figsize = [15,15])

ax = fig.gca(projection = "3d")
ax.view_init(elev=None, azimuth=None) #you can set elevation/azimuth with different values

ax.plot(points_3D[0], points_3D[1], points_3D[2], 'o')

plt.draw()
plt.show()
```

Now, we perform 3D projection for all the 3D points in [house.p3d](#) onto an image using the camera matrix [P](#) defined above.

```
#projection
points_2D = project(P, points_3D)

#plot projection
from matplotlib import pyplot as plt
plt.plot(points_2D[0], points_2D[1], 'k.')
plt.show()
```

You can change the camera matrix [P](#) and observe the results of 3D projection accordingly.

2. Camera pose estimation

In this section, we learn how to estimate the camera matrix [P](#) of a pin-hole camera given a set of 3D points and their projected 2D locations. Suppose that we are given [points_3D](#) as a set of 3D points and [points_2D](#) as their 2D locations projected on an image. Suppose that [P](#) is unknown and we want to estimate it.

In Eq (7) in slide 14, Lecture 6 slides, we have learnt how to build a matrix [A](#) that can be used for estimation of [P](#). We have also learnt that we need at least 6 points to build the matrix [A](#) (see slide 17, Lecture 6 slides). However, [points_3D](#) contains more than 6 points. We can check this using the following commands.

```
print(points_2D.shape)
print(points_3D.shape)
```

To build the matrix [A](#), we simply sample [points_3D](#) by taking the first 6 points from [points_3D](#) and their corresponding 2D locations from [points_2D](#). The sampled 3D points and corresponding 2D locations are stored in [points_3D_sampled](#) and [points_2D_sampled](#) respectively.

```
n_points = 6
points_3D_sampled = points_3D[:, :n_points]
points_2D_sampled = points_2D[:, :n_points]
```

We now build the matrix A from `points_3D_sampled` and `points_2D_sampled` as follows.

```
A = np.zeros((2*n_points, 12), np.float32)
for i in range(n_points):
    A[2*i,:4] = points_3D_sampled[:,i].T
    A[2*i,8:12] = -points_2D_sampled[0,i] * points_3D_sampled[:,i].T
    A[2*i+1,4:8] = points_3D_sampled[:,i].T
    A[2*i+1,8:12] = -points_2D_sampled[1,i] * points_3D_sampled[:,i].T
```

Given A , our problem is to find p such that $Ap = 0$. However, we have learnt that solving such a linear system would result in trivial solution. We can double check this simply as,

```
from scipy import linalg
p = linalg.solve(A, np.zeros((12, 1), np.float32))
print(p)
```

As shown, all the elements of p are zero, i.e., trivial solution. To address this issue, instead of directly solving $Ap = 0$, we minimise $\|Ap\|$ s.t. $\|p\| = 1$ using SVD (singular value decomposition) technique. In particular, we first decompose A as,

```
U, S, V = linalg.svd(A)
```

Results of the above command include a matrix U , a vector S , and a matrix V satisfying $U \cdot \text{diag}(S) \cdot V^T = A$. Note that S is a sorted array of singular values with decreasing order. S contains 12 values and thus the $S[11]$ is the smallest singular value. We can verify this using the following command.

```
minS = np.min(S)
conditon = (S == minS)
minID = np.where(conditon)
print('index of the smallest singular value is: ', minID[0])
```

We finally compute P_{hat} (an estimate of P) by performing:

- Taking the row of V (i.e., column of V^T) that corresponds to the smallest singular value (i.e., the last row)
- Express this column in a (3×4) matrix form
- Divide all elements of the matrix by the smallest singular value

```
P_hat = V[minID[0],:].reshape(3, 4) / minS
```

Now, we print both P and P_{hat} .

```
print(P)
print(P_hat)
```

P_{hat} does not seem to be close to P ! That's ok! In fact, two different camera matrices may produce the same projected image due to the conversion from homogeneous to inhomogeneous. We now perform another test, that is, to take a 3D point, project this 3D point to 2D using both P and P_{hat} and then compare the projection results. To do so, we simply take the first point in `points_3D_sampled` and project it to 2D using P_{hat} .

```
x_P_hat = project(P_hat, points_3D_sampled[:, 0])
print(x_P_hat)
```

Now, we print the 2D location of that point from `points_2D_sampled` obtained by using P .

```
x_P = points_2D_sampled[:,0]
print(x_P)
```

For a quantitative evaluation, we measure the distance between x_P and x_{P_hat} for all points in `points_3D`. The smaller the distance is, the more accurate the estimation (of P) is.

```
x_P = points_2D
x_P_hat = project(P_hat, points_3D)

dist = 0
for i in range(x_P.shape[1]):
    dist += np.linalg.norm(x_P[:,i] - x_P_hat[:,i])
dist /= x_P.shape[1]
print(dist)
```

Your task is to vary `n_points` in the range [10%, 20%, ..., 100%] of the number of points in `points_3D` and measure the corresponding `dist`. What is the best value for `n_points`?

3. 3D reconstruction

In this section, we will experiment the structure-from-motion (sfm) approach for 3D reconstruction from two views with assumption that the intrinsic calibration matrix K is given. We will use the two images `temple2A.png` and `temple2B.png` supplied in OnTrack. These images come from the paper “A comparison and evaluation of multi-view stereo reconstruction algorithms” by Seitz et al. in CVPR 2006, and can be downloaded at <https://vision.middlebury.edu/mview/>. The images capture a scene (a temple) at two different viewpoints.

We also use some code from this book (with minor modifications to adapt with Python 3) including `homography.py`, `ransac.py`, and `sfm.py`. The code is provided in OnTrack and can also be found at <https://github.com/jesolem/PCV>. You are also recommended to refer to <https://www.opensfm.org/> for further study on the sfm approach.

You first need to download the supplied `homography.py`, `ransac.py`, and `sfm.py` into your working directory and then import them to your notebook. You can read a description of the RANSAC algorithm (implemented in `ransac.py`) at https://en.wikipedia.org/wiki/Random_sample_consensus.

```
import homography
import sfm
import ransac
```

You also need to download the two image files `temple2A.png` and `temple2B.png` into your working directory. You then read the images from those files and extract SIFT keypoints and descriptors from them. SIFT keypoints and descriptors on `temple2A.png` are stored in `kp1` and `des1`, and those on `temple2B.png` are stored in `kp2` and `des2`, respectively.

```
import cv2 as cv

sift = cv.xfeatures2d.SIFT_create()

img1 = cv.imread('temple2A.png')
img1_gray = cv.cvtColor(img1, cv.COLOR_BGR2GRAY)
kp1, des1 = sift.detectAndCompute(img1_gray, None)

img2 = cv.imread('temple2B.png')
img2_gray = cv.cvtColor(img2, cv.COLOR_BGR2GRAY)
kp2, des2 = sift.detectAndCompute(img2_gray, None)
```

```

img1_kp = img1.copy()
img1_kp = cv.drawKeypoints(img1, kp1, img1_kp)

print("Number of detected keypoints in img1: %d" % (len(kp1)))

img2_kp = img2.copy()
img2_kp = cv.drawKeypoints(img2, kp2, img2_kp)

print("Number of detected keypoints in img2: %d" % (len(kp2)))

img1_2_kp = np.hstack((img1_kp, img2_kp))

plt.figure(figsize = (20, 10))
plt.imshow(img1_2_kp[:, :, ::-1])
plt.axis('off')

```

We match descriptors in [des1](#) and those in [des2](#) using [cv.BFMatcher](#) as follows.

```

bf = cv.BFMatcher(crossCheck = True) #crossCheck=True: find consistent matches
matches = bf.match(des1, des2)
matches = sorted(matches, key = lambda x:x.distance)
print("Number of consistent matches: %d" % len(matches))

```

The best 20 matches are visualised as follows.

```

img1_2_matches = cv.drawMatches(img1, kp1, img2, kp2,
                                matches[:20],
                                None,
                                flags = cv.DrawMatchesFlags_NOT_DRAW_SINGLE_POINTS)

plt.figure(figsize = (20, 10))
plt.imshow(img1_2_matches[:, :, ::-1])
plt.axis('off')

```

Next, we identify the corresponding matching keypoints and store them in arrays for later use. Note that high number of matches would lead to high computational complexity for the reconstruction. Therefore, in some cases, one can define an upper bound for the number of matches, e.g., 1000.

```

n_matches = min(len(matches), 1000)

kp1_array = np.zeros((2, n_matches), np.float32)
for i in range(n_matches):
    kp1_array[0][i] = kp1[matches[i].queryIdx].pt[0]
    kp1_array[1][i] = kp1[matches[i].queryIdx].pt[1]

kp2_array = np.zeros((2, n_matches), np.float32)
for i in range(n_matches):
    kp2_array[0][i] = kp2[matches[i].trainIdx].pt[0]
    kp2_array[1][i] = kp2[matches[i].trainIdx].pt[1]

```

We convert all the keypoints into homogeneous coordinates.

```

x1 = homography.make_homog(kp1_array)
x2 = homography.make_homog(kp2_array)

```

Suppose that the intrinsic calibration matrix K is given, and we fix $P1=[I|0]$, we need to find $P2$ (see slides 49-50, Lecture 6 slides). For example, we define:

```
K = np.array([[2394,0,932], [0,2398,628], [0,0,1]])
P1 = np.array([[1,0,0,0], [0,1,0,0], [0,0,1,0]])
```

We normalise **x1** and **x2** using K^{-1} .

```
x1n = np.dot(linalg.inv(K), x1)
x2n = np.dot(linalg.inv(K), x2)
```

We estimate the essential matrix **E**. **Note:** this step takes time.

```
#estimate E with RANSAC
model = sfm.RansacModel()
E, inliers = sfm.F_from_ransac(x1n, x2n, model)
```

We then compute **P2**.

```
#compute camera matrices (P2 will be list of four solutions)
P2_all = sfm.compute_P_from_essential(E)

#pick the solution with points in front of cameras
ind = 0
maxres = 0
for i in range(4):
    #triangulate inliers and compute depth for each camera
    X = sfm.triangulate(x1n[:, inliers], x2n[:, inliers], P1, P2_all[i])
    d1 = np.dot(P1, X)[2]
    d2 = np.dot(P2_all[i], X)[2]
    s = sum(d1 > 0) + sum(d2 > 0)
    if s > maxres:
        maxres = s
        ind = i
        infront = (d1 > 0) & (d2 > 0)
P2 = P2_all[ind]
```

Finally, we reconstruct 3D points using the triangulation algorithm (see Eq (20) in slide 52, Lecture 6 slides).

```
#triangulate inliers and remove points not in front of both cameras
X = sfm.triangulate(x1n[:, inliers], x2n[:, inliers], P1, P2)
X = X[:, infront]
```

After applying the triangulation algorithm and removing points not in front of both cameras, we obtain a set of reconstructed 3D points, named **X**. This set contains less than 1000 points. We can check this by performing.

```
print(len(X[0]))
```

We visualise **X** using the following code.

```
#3D plot
fig = plt.figure(figsize = [20,20])

ax = fig.gca(projection = "3d")
ax.view_init(elev = -80, azimuth = -90)

ax.plot(X[0], X[1], X[2], 'o')

plt.draw()
plt.show()
```

Your task is to visualise the original points (computed using SIFT) on [img2](#) and the image points (computed by projecting X with camera matrix P_2) on [img2](#).

Hint: After projecting X to x by using “project” method, you need to convert x into image coordinates by multiply K with x .

Submission instructions

1. Perform tasks required in Section 2 and 3.
2. Complete the supplied answer sheet with required results
3. Submit the answer sheet (.pdf) and code (.py) to OnTrack.